

Rollercoaster

Gao

2023-10-30

Import libraries or something

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.2      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2    3.4.3      v tibble    3.2.1
## v lubridate  1.9.2      v tidyr     1.3.0
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(tidymodels)
```

```
## -- Attaching packages ----- tidymodels 1.1.1 --
## v broom      1.0.5      v rsample    1.2.0
## v dials      1.2.0      v tune       1.1.2
## v infer      1.0.5      v workflows  1.1.3
## v modeldata  1.2.0      v workflowsets 1.0.1
## v parsnip    1.1.1      v yardstick  1.2.0
## v recipes    1.0.8
## -- Conflicts ----- tidymodels_conflicts() --
## x scales::discard() masks purrr::discard()
## x dplyr::filter()   masks stats::filter()
## x recipes::fixed() masks stringr::fixed()
## x dplyr::lag()      masks stats::lag()
## x yardstick::spec() masks readr::spec()
## x recipes::step()   masks stats::step()
## * Learn how to get started at https://www.tidymodels.org/start/
```

```
library(ggforce)
library(yardstick)
library(car)
```

```
## Loading required package: carData
##
```

```
## Attaching package: 'car'
##
## The following object is masked from 'package:dplyr':
##
##   recode
##
## The following object is masked from 'package:purrr':
##
##   some
```

```
library(moments)
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg   ggplot2
```

```
library(psych)
```

```
##
## Attaching package: 'psych'
##
## The following object is masked from 'package:car':
##
##   logit
##
## The following objects are masked from 'package:scales':
##
##   alpha, rescale
##
## The following objects are masked from 'package:ggplot2':
##
##   %+%, alpha
```

```
library(fastDummies)
```

```
## Thank you for using fastDummies!
## To acknowledge our work, please cite the package:
## Kaplan, J. & Schlegel, B. (2023). fastDummies: Fast Creation of Dummy (Binary) Columns and Rows from
```

Import the data

```
rollercoasters <- read_csv("rollercoasters.csv") %>% as_tibble()
```

```
## Rows: 101 Columns: 9
## -- Column specification -----
## Delimiter: ","
## chr (3): Name, Park, Track
## dbl (6): Speed, Height, Drop, Length, Duration, Inversions
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

Find correlation between Speed and all numerical values

```
cor(rollercoasters$Speed, select_if(rollercoasters, is.numeric))
```

```
##      Speed      Height      Drop      Length      Duration      Inversions
## [1,]      1 0.9170502 0.9543598 0.4562342 0.1704023 -0.4037699
```


Find the regression model

```
model <- lm(Speed ~ Height + Drop + Length + Duration + Inversions, data = rollercoasters)
summary(model)
```

```
##
## Call:
## lm(formula = Speed ~ Height + Drop + Length + Duration + Inversions,
##     data = rollercoasters)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.8097 -2.4500 -0.1268  1.6911 17.7084
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 35.7956667  1.7295008  20.697  < 2e-16 ***
## Height      0.0406960  0.0161666   2.517  0.0135 *
## Drop        0.1468456  0.0178691   8.218 1.04e-12 ***
## Length      0.0013097  0.0005315   2.464  0.0155 *
## Duration    -0.0293653  0.0159234  -1.844  0.0683 .
## Inversions  -0.2885555  1.0018872  -0.288  0.7740
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.034 on 95 degrees of freedom
## Multiple R-squared:  0.9228, Adjusted R-squared:  0.9187
## F-statistic: 227 on 5 and 95 DF, p-value: < 2.2e-16
```

Construct the residual plot

```
plot(fitted(model), resid(model))
```



Draw outlier boundaries


```
standard_error <- sqrt(deviance(model)/df.residual(model))  
standard_error
```

```
## [1] 4.034104
```

```
2*standard_error
```

```
## [1] 8.068208
```

```
plot(fitted(model),resid(model))  
abline(h=2*standard_error, col = "red")  
abline(h=-2*standard_error, col = "red")
```



2x standard error of mean: 8.068208 There are 5 outliers that are outside of 2 standard errors of the mean.
(Coasters 19, 2, 16, 43, 3) Display residual values

```
residuals <- resid(model)
sort(residuals)
```

```
##          19          2          16          17          24          9
## -9.80968692 -9.43582323 -9.14564310 -6.88655153 -6.57873045 -5.03502709
##          6          29          36          37          38          39
## -4.46403599 -3.87797696 -3.81426290 -3.81426290 -3.81426290 -3.81426290
##          40          41          59          80          11          67
## -3.81426290 -3.81426290 -3.77848582 -3.65532739 -3.53095820 -3.25728330
##          45          99          13          14          15          12
## -2.98118649 -2.79099941 -2.57140433 -2.57140433 -2.57140433 -2.56733473
##          84          42          101          89          61          20
## -2.56308696 -2.45003243 -2.39628114 -2.30258696 -2.23405392 -1.98531321
##          86          35          87          92          79          60
## -1.92910179 -1.81035928 -1.71792525 -1.65979658 -1.65366258 -1.64575621
##          82          75          76          28          57          74
## -1.54901403 -1.29024392 -1.08676379 -1.06388347 -0.95705821 -0.86666831
##          68          81          10          27          52          97
## -0.86365343 -0.84069790 -0.61621282 -0.57933503 -0.54028131 -0.34150390
##          64          65          1          70          96          100
## -0.26324045 -0.20740556 -0.12677727 -0.07609965  0.35354490  0.38856557
##          21          22          66          62          69          90
##  0.40898972  0.40898972  0.46269561  0.52785803  0.74986959  0.79463263
```

```
##          85          58          56          8          71          49
## 0.99778771 1.12009791 1.22804365 1.26706370 1.26723237 1.32461850
##          4          72          73          48          78          77
## 1.41862094 1.51296753 1.52178776 1.55878547 1.58726310 1.64038326
##          31          32          33          34          23          47
## 1.69113548 1.69113548 1.69113548 1.69113548 1.72204787 1.82570548
##          44          83          25          26          46          63
## 1.89005657 1.93160418 1.93289874 1.93289874 2.81728320 2.96707347
##          55          7          50          88          91          30
## 2.97044418 3.05763232 3.08368177 3.24381327 3.25561435 3.51387496
##          94          18          5          93          51          98
## 3.72883602 3.94274100 4.76066630 5.58758210 5.69255297 6.04419386
##          95          53          54          43          3
## 7.12571531 7.55162527 7.70435850 10.71597688 17.70841942
```

Cook's Distance

Jackknife

```
jac <- qt(df= 101 - 5 - 2, 0.95)
jac
```

```
## [1] 1.661226
```

```
sort(jackknife <- rstudent(model))
```

```
##          19          2          16          17          24          9
## -2.57052918 -2.52245096 -2.43993818 -1.77919507 -1.69086432 -1.28471954
##          6          29          36          37          38          39
## -1.14891029 -0.98813891 -0.95610805 -0.95610805 -0.95610805 -0.95610805
##          40          41          59          80          11          67
## -0.95610805 -0.95610805 -0.94932547 -0.93775067 -0.89029940 -0.84329770
##          45          99          13          14          15          12
## -0.76252985 -0.74778427 -0.65357322 -0.65357322 -0.65357322 -0.65253416
##          84          101          42          89          61          20
## -0.64584669 -0.62994342 -0.61819026 -0.58790073 -0.56226222 -0.50327540
##          86          92          35          87          79          60
## -0.49250273 -0.47683253 -0.45935716 -0.43919044 -0.41715334 -0.41233994
##          82          75          76          28          57          68
## -0.38895493 -0.32296354 -0.27207075 -0.26759471 -0.24004070 -0.22607518
##          74          81          10          27          52          97
## -0.21685351 -0.21344133 -0.15479042 -0.14492978 -0.13603176 -0.08524739
##          64          65          1          70          96          100
## -0.06597490 -0.05237858 -0.03491871 -0.01909104 0.09003219 0.09702678
##          21          22          66          62          69          90
## 0.10286584 0.10286584 0.11696652 0.13591656 0.18817305 0.20759802
##          85          58          56          71          8          49
## 0.25489372 0.28053277 0.30776545 0.31822233 0.32301855 0.34188350
##          4          73          72          48          78          77
## 0.35999469 0.38084818 0.38481067 0.39147570 0.40459346 0.41135923
##          31          32          33          34          23          47
## 0.42325485 0.42325485 0.42325485 0.42325485 0.43209228 0.45753479
##          25          26          44          83          46          63
```

```
## 0.48380471 0.48380471 0.48566289 0.48916676 0.71809364 0.74520295
##          55          50           7           88           91           30
## 0.74607235 0.77502879 0.81556456 0.83499983 0.85298574 0.88162447
##          18          94           5           51           93           98
## 1.00901671 1.11794151 1.21248390 1.46132284 1.60427800 1.61927363
##          95          53          54          43           3
## 1.86068145 1.94145663 1.96255671 2.79804194 5.32771209
```

Jackknife = 1.661226 Outliers: 19, 2, 16, 17, 24, 95, 53, 54, 43, 3

Leverage

```
lev <- 2 * (1+5)/101
lev
```

```
## [1] 0.1188119
```


```
sort(hatvalues(model))
```

```
##          36          37          38          39          40          41          97
## 0.02294170 0.02294170 0.02294170 0.02294170 0.02294170 0.02294170 0.02417290
##          54          100          30          25          26          31          32
## 0.02460875 0.02478805 0.02615072 0.02709997 0.02709997 0.02749932 0.02749932
##          33          34          59          73          27          74          75
## 0.02749932 0.02749932 0.02756770 0.02773617 0.02825684 0.02837150 0.02853224
##          76          60          47          58          63          55          56
## 0.02913558 0.02967965 0.02973714 0.02989593 0.03043768 0.03048370 0.03097406
##          50          77          64          23          57          70          43
## 0.03132095 0.03141134 0.03199412 0.03237011 0.03287375 0.03390946 0.03392911
##          69          82          48          71          11          10          61
## 0.03410409 0.03412559 0.03443687 0.03477419 0.03557196 0.03618410 0.03688585
##          84          28          21          22          52          42          79
## 0.03816450 0.03822410 0.03873882 0.03873882 0.04069913 0.04110397 0.04277133
##          53          65           5          66          83           9          20
## 0.04322453 0.04664171 0.04800815 0.04842953 0.04953390 0.04970929 0.05130761
##          24          19          35          29           4          12          13
## 0.05160556 0.05228370 0.05351776 0.05382278 0.05452730 0.05456554 0.05456555
##          14          15          81          51          17          72          46
## 0.05456555 0.05456555 0.05627943 0.05640049 0.05843355 0.05863353 0.05901029
##          18          96          78           8          89          86          45
## 0.06159903 0.06234935 0.06259798 0.06344368 0.06388735 0.06476160 0.06491126
##          80          85          87           6          95          88          44
## 0.06753644 0.06767307 0.06781281 0.06921494 0.07544861 0.07560322 0.07683728
##          62          67          49           2          16          91          90
## 0.08275575 0.08602860 0.08614950 0.09161704 0.09165624 0.10743247 0.10875816
##          68          101           3          98           7          99           1
## 0.11219125 0.11648868 0.12544559 0.12924741 0.13935076 0.14797409 0.19853929
##          93          92          94
## 0.24224338 0.26152245 0.31458327
```

leverage = 0.1188119 Outliers: 3, 98, 7, 99, 1, 93, 92, 94

QQPlot

```
qqnorm(resid(model))
```



```
model <- lm(Speed ~ Height + Drop + Length + Duration + Inversions, data = rollercoasters)
model
```

```
##
## Call:
## lm(formula = Speed ~ Height + Drop + Length + Duration + Inversions,
##     data = rollercoasters)
##
## Coefficients:
## (Intercept)      Height          Drop          Length      Duration      Inversions
##   35.79567      0.04070      0.14685      0.00131     -0.02937     -0.28856
```

Speed-hat = 0.0407 * Height + 0.14685 * Drop + 0.00131 * Length + -0.2937 * Duration + -0.28856 * Inversions + 35.79567

```
rollercoasters <- dummy_cols(rollercoasters, select_columns = "Track", remove_first_dummy = TRUE)
rollercoasters
```

```
## # A tibble: 101 x 10
##   Name      Park Track Speed Height Drop Length Duration Inversions Track_Wood
##   <chr>    <chr> <chr> <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>    <int>
## 1 Tower o~ Drea~ Steel   100   377.  328.  1235      28         0         0
```



```
## 2 Canyon ~ Adve~ Steel 41 94 66 2423 60 1 0
## 3 Xcelera~ Knot~ Steel 82 205 130 2202 62 0 0
## 4 Afterbu~ Fun ~ Steel 45 56 47 635 66 1 0
## 5 Silver ~ Fron~ Steel 55 83 75 1942 75 1 0
## 6 New Mex~ Clif~ Wood 47 80 75 2750 75 0 1
## 7 Outlaw ~ Silv~ Wood 68 107 162 2937 87 1 1
## 8 Thunder~ Kenn~ Wood 55 70 95 2887 90 0 1
## 9 Inverti~ Para~ Steel 55 138 138 985 90 1 0
## 10 Freesty~ Cava~ Steel 51 88 84 2210 92 1 0
## # i 91 more rows
```

```
model2 <- lm(Speed ~ Height + Drop + Length + Duration + Inversions + Track_Wood, data = rollercoasters)
summary(model2)
```

```
##
## Call:
## lm(formula = Speed ~ Height + Drop + Length + Duration + Inversions +
##     Track_Wood, data = rollercoasters)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.1249  -2.3337  -0.2069   1.8148  18.1630
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 34.8044115  2.1521406  16.172  < 2e-16 ***
## Height      0.0423121  0.0163335   2.591  0.0111 *
## Drop        0.1490159  0.0181231   8.222 1.09e-12 ***
## Length      0.0011928  0.0005534   2.155  0.0337 *
## Duration    -0.0264514  0.0163916  -1.614  0.1099
## Inversions   0.1339013  1.1417891   0.117  0.9069
## Track_Wood   1.0172394  1.3093781   0.777  0.4392
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.043 on 94 degrees of freedom
## Multiple R-squared:  0.9232, Adjusted R-squared:  0.9183
## F-statistic: 188.5 on 6 and 94 DF,  p-value: < 2.2e-16
```

```
model2 %>% tidy()
```

```
## # A tibble: 7 x 5
##   term      estimate std.error statistic  p.value
##   <chr>      <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept) 34.8       2.15      16.2 6.68e-29
## 2 Height      0.0423    0.0163     2.59 1.11e- 2
## 3 Drop        0.149     0.0181     8.22 1.09e-12
## 4 Length      0.00119   0.000553   2.16 3.37e- 2
## 5 Duration    -0.0265    0.0164    -1.61 1.10e- 1
## 6 Inversions   0.134     1.14       0.117 9.07e- 1
## 7 Track_Wood   1.02      1.31       0.777 4.39e- 1
```